

CS 411: Artificial Intelligence I

Programming Assignment 2

Due: Sunday of Week 5, 11:59pm

This portion of the assignment may be completed individually or in groups of 2.

In the programming portion of this assignment we will use a number of python packages to explore some of the tools we have learned about in class. Depending on your setup, this may require some installation. The following instructions were tested on a fresh installation of python 3.6.4.

On Windows, before installing the packages, if you do not already have a recent version you will need to install the Visual Studio Standalone C++ Build Tools <https://visualstudio.microsoft.com/downloads/#build-tools-for-visual-studio-2019>, or Visual Studio with C++ Tools <https://visualstudio.microsoft.com/vs/features/cplusplus/>. For more detailed help with this, see <https://wiki.python.org/moin/WindowsCompilers>. If this large download is problematic for you, see the comment on Google Colab below. Mac and Linux typically come with suitable compilers already installed.

Then download the optimization.zip file from Blackboard and unzip it so you have access to the requirements.txt file. Go to the homework directory and run the following command to install the packages and their dependencies:

```
pip install -r requirements.txt
```

If your path isn't set up correctly for pip you may instead need to do

```
python -m pip install -r requirements.txt
```

Be sure to check the output of each to verify the installation occurred without error. For example, you may get an error if you failed to install the above compiler first. If you have issues installing cvxpy, see the instructions at <http://www.cvxpy.org/install/index.html>.

Once you have successfully run these, you will have installed the packages we will be using for convex programming (cvxpy), simulated annealing (simanneal), and constraint satisfaction (python-constraint).

If you have problems in setting up your homework programming environment or do not want to install C++ compiler, you can also use Google Colab as your working environment (See separate instruction for details). If you run into issues getting set up, check existing threads and start a new one if needed on Piazza and we will do our best to help.

When running the autograder, the fifth problem will take some time to run, even if you haven't implemented it yet. You may wish to run individual questions using the following syntax, which runs just question 1.

```
python autograder.py -q 1
```

To help you get started with some of the packages we have included some demo files. Feel free to run, read, and edit these to check out the examples. The only file for the actual assignment you need to edit is optimization.py.

1. **CSP: Sudoku (5 points).** We have provided most of a CSP implementation of Sudoku. You need to implement `cstAdd`, which adds the constraints. It takes a problem object (problem), a matrix of variable names (grid), a list of legal values (domains), and the side length of the inner squares (psize, which is 3 in an ordinary sudoku and 2 in the smaller version we provide as the easier test case). That is, the Sudoku puzzle consists of a psize by psize grid of psize by psize squares for psize⁴ cells in total
2. **Linear Programming: Fractional Knapsack (5 points).**

In the fractional knapsack problem you have a knapsack with a fixed weight capacity and want to fill it with valuable items so that we maximize the total value in it while ensuring the weight does not exceed the capacity c . Fractions of items are allowed. Implement the function `fractionalKnapsack` to use an LP to solve the fractional knapsack problem with 3 items of weights 5, 3, and 1 and values 2, 3, and 1 respectively for varying knapsack capacities. There is only one unit of each item available.

3. **Integer Programming: Sudoku (5 points).** For our final Sudoku, we have provided most of an IP implementation. Again, you just need to implement the constraints. Note however, unlike in the CSP version, we have not already “prefilled” the squares for you. You’ll need to add those constraints yourself.
4. **Local Search: TSP (5 points).** We have provided most of a simulated annealing implementation of the famous traveling salesman problem, where you seek to visit a list of cities while minimizing the total distance traveled. You need to implement `move` and `energy`. The former is the operation for finding nearby candidate solutions while the latter evaluates how good the current candidate solution is. For this particular problem, `move` should generate a random local move without regard for whether it is beneficial. Similarly, to receive credit `energy` should calculate the total euclidean distance of the current candidate tour. There is a `distance` function you may wish to implement to help with this.
5. **Local Search: Sudoku (5 points).** Now we have the skeleton of a simulated annealing implementation of Sudoku. You need to design the `move` and `energy` functions and will receive credit based on how many of 10 runs succeed in finding a correct answer: to achieve k points $2k - 1$ runs need to pass.

Your code in the file `optimization.py` should be submitted via gradescope for evaluation. **It must be your own (or you and your partner’s) code and should not be copied from any other source.** We will check for similarity to other submissions and existing resources available on the web for any cheating. Abuse of the autograder, e.g. by hardcoding solutions, is also considered cheating.

Detailed submission instructions:

1. Please submit *only* the listed files and (optionally) a README file via Gradescope
2. You can submit the files individually or submit a zip of them all, but if you zip them make sure they are in the root of the zip. (That is, make sure to zip the files themselves, not the directory that contains them.)
3. The README is to help us give you partial credit if there are any issues. Things you could comment on include anything special we need to do to run your code (e.g. if you used a non-standard package), any bugs we need to work around, and if there is a partially complete implementation which contains material you commented out for one of the questions so that the overall code will run.
4. If you worked with a partner, make sure to do it as a group submission on gradescope so that all of you get credit.

Grading standards:

- Full Credit (1): ≥ 20 points on autograder
- Partial Credit (0.5): 10 – 19 points on autograder
- Some Credit (0.25): 5 – 9 points on autograder
- Insufficient evidence (0): ≤ 4 points on autograder